

DATA Cleaning

Now that we have a better understanding of our data set, let's try to get a cleaner data set.

Table of Contents

1. [Step 1: Dealing with NaN values](#)
2. [Step 2: Converting data types](#)
3. [Step 3: Categorical variables](#)
4. [Step 4: Save our cleaned dataset for reusability](#)
5. [Conclusion](#)

Step 1: Dealing with NaN values

We noticed during our EDA that specific data variables contain NaN values :

- yr_renovated
- sqft_basement
- view
- waterfront

```
In [1]: import pandas as pd
data = pd.read_csv('kc_house_data.csv')
data.head()
```

Out[1]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	NaN
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	0.0
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	0.0
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	0.0
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	0.0

5 rows × 21 columns

Year Renovated

```
In [2]: # Fill NaN Year Renovated with Yr Built
cleaned_data = data.copy()
cleaned_data['yr_renovated'] = cleaned_data['yr_renovated'].fillna(cleaned_data['yr_built'])
```

```
In [3]: # Fill 0 Year renovated with Yr Built
cleaned_data.loc[cleaned_data['yr_renovated'] == 0, 'yr_renovated'] = cleaned_data['yr_built']
cleaned_data.head()
```

Out[3]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	NaN
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	0.0
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	0.0
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	0.0
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	0.0

5 rows × 21 columns

```
In [4]: cleaned_data.shape
```

Out[4]: (21597, 21)

Sqft Basement

```
In [5]: # Remove non numerical values and set them to null
print('sqft_basement - Count non numeric values : ', len(cleaned_data.loc[cleaned_data['sqft_basement'] == '?', 'sqft_basement']))
cleaned_data.head()
```

sqft_basement - Count non numeric values : 454

Out[5]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	NaN
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	0.0
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	0.0
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	0.0
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	0.0

5 rows × 21 columns

View

```
In [6]: # Check different values for view
print(cleaned_data['view'].value_counts())
print('view - Total of missing values : ', cleaned_data['view'].isna().sum())

0.0    19422
2.0     957
3.0     508
1.0     330
4.0     317
Name: view, dtype: int64
view - Total of missing values : 63
```

```
In [7]: # Set missing values to null
cleaned_data['view'].fillna(0, inplace=True)
cleaned_data.isna().sum()
```

```
Out[7]: id                0
date                    0
price                   0
bedrooms                0
bathrooms               0
sqft_living              0
sqft_lot                 0
floors                  0
waterfront              2376
view                    0
condition               0
grade                   0
sqft_above              0
sqft_basement           0
yr_built                0
yr_renovated            0
zipcode                 0
lat                     0
long                    0
sqft_living15           0
sqft_lot15              0
dtype: int64
```

Waterfront

```
In [8]: # Check different values for waterfront
print(cleaned_data['waterfront'].value_counts())
print('Waterfront - Total of missing values : ', cleaned_data['waterfront'].isna().sum())

0.0    19075
1.0     146
Name: waterfront, dtype: int64
Waterfront - Total of missing values : 2376
```

```
In [9]: # Set missing values to null
cleaned_data['waterfront'].fillna(0, inplace=True)
cleaned_data.isna().sum()
```

```
Out[9]: id                0
        date              0
        price             0
        bedrooms          0
        bathrooms         0
        sqft_living       0
        sqft_lot          0
        floors            0
        waterfront       0
        view              0
        condition         0
        grade             0
        sqft_above        0
        sqft_basement     0
        yr_built          0
        yr_renovated      0
        zipcode           0
        lat               0
        long              0
        ...
```

Conclusion During the NaN/Null values analysis we made the choice of keeping as max as possible entries given values. We did it by assigning default value to the most commonly value in dataset (or mode value). Next step, now that we have cleaned our missing data, let's fix the datatypes.

Step 2: Converting data types

As shown previously, we have deal with NaN & null values. It is time to convert our dataset with correct datatype.

Looking back to our EDA, we need to convert the following data to the correct datatype:

- date to datetime
- waterfront to bool
- yr_built to datetime
- yr_renovated to datetime

Transforming date, yr_built, yr_renovated to datetime

```
In [10]: # Transform date to datetime using pd.to_datetime()
cleaned_data['date'] = pd.to_datetime(cleaned_data['date'], format='%m/%d/%
cleaned_data.dtypes
```

```
Out[10]: id                int64
date                datetime64[ns]
price                float64
bedrooms            int64
bathrooms           float64
sqft_living         int64
sqft_lot            int64
floors              float64
waterfront          float64
view                float64
condition           int64
grade               int64
sqft_above          int64
sqft_basement       object
yr_built            int64
yr_renovated         float64
zipcode             int64
lat                 float64
long                float64
sqft_living15       int64
sqft_lot15          int64
dtype: object
```

```
In [11]: # Transform yr_built to datetime
cleaned_data['yr_built'] = pd.to_datetime(cleaned_data['yr_built'], format=
cleaned_data.dtypes
```

```
Out[11]: id                int64
date                datetime64[ns]
price                float64
bedrooms            int64
bathrooms           float64
sqft_living         int64
sqft_lot            int64
floors              float64
waterfront          float64
view                float64
condition           int64
grade               int64
sqft_above          int64
sqft_basement       object
yr_built            datetime64[ns]
yr_renovated         float64
zipcode             int64
lat                 float64
long                float64
sqft_living15       int64
sqft_lot15          int64
dtype: object
```

```
In [12]: # Transform yr_renovated to datetime in 2 steps. First converting into int
cleaned_data['yr_renovated'] = cleaned_data['yr_renovated'].apply(lambda x:
cleaned_data['yr_renovated'] = pd.to_datetime(cleaned_data['yr_renovated'],
cleaned_data.dtypes
```

```
Out[12]: id                                int64
date                                datetime64[ns]
price                                float64
bedrooms                            int64
bathrooms                           float64
sqft_living                          int64
sqft_lot                             int64
floors                              float64
waterfront                          float64
view                                float64
condition                           int64
grade                               int64
sqft_above                          int64
sqft_basement                        object
yr_built                            datetime64[ns]
yr_renovated                        datetime64[ns]
zipcode                             int64
lat                                 float64
long                                float64

```

Transforming waterfront to boolean

```
In [13]: # Check waterfront values
cleaned_data['waterfront'] = cleaned_data['waterfront'].astype(bool)
cleaned_data.dtypes
```

```
Out[13]: id                                int64
date                                datetime64[ns]
price                                float64
bedrooms                            int64
bathrooms                           float64
sqft_living                          int64
sqft_lot                             int64
floors                              float64
waterfront                          bool
view                                float64
condition                           int64
grade                               int64
sqft_above                          int64
sqft_basement                        object
yr_built                            datetime64[ns]
yr_renovated                        datetime64[ns]
zipcode                             int64
lat                                 float64
long                                float64

```

Step 3: Categorical variables

Do we have categorical variables ? We have some insights from our EDA about which are the right candidates but let's dig in.

```
In [14]: # Spotting Categorical Variables
cleaned_data.describe()
```

Out[14]:

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	
count	2.159700e+04	2.159700e+04	21597.000000	21597.000000	21597.000000	2.159700e+04	21597
mean	4.580474e+09	5.402966e+05	3.373200	2.115826	2080.321850	1.509941e+04	1.
std	2.876736e+09	3.673681e+05	0.926299	0.768984	918.106125	4.141264e+04	0.
min	1.000102e+06	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+02	1.
25%	2.123049e+09	3.220000e+05	3.000000	1.750000	1430.000000	5.040000e+03	1.
50%	3.904930e+09	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03	1.
75%	7.308900e+09	6.450000e+05	4.000000	2.500000	2550.000000	1.068500e+04	2.
max	9.900000e+09	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3.

At first glance, we can already spot some categorical variables : condition, grade, bedrooms, bathrooms, floors and maybe view / zipcode

```
In [15]: # Check uniqueness
categorical_variables = ['waterfront', 'condition', 'grade', 'bedrooms', 'b
for categorical_variable in categorical_variables:
    print('{} unique values : '.format(categorical_variable), cleaned_data
```

```
waterfront unique values : 2
condition unique values : 5
grade unique values : 11
bedrooms unique values : 12
bathrooms unique values : 29
floors unique values : 6
view unique values : 5
zipcode unique values : 70
```

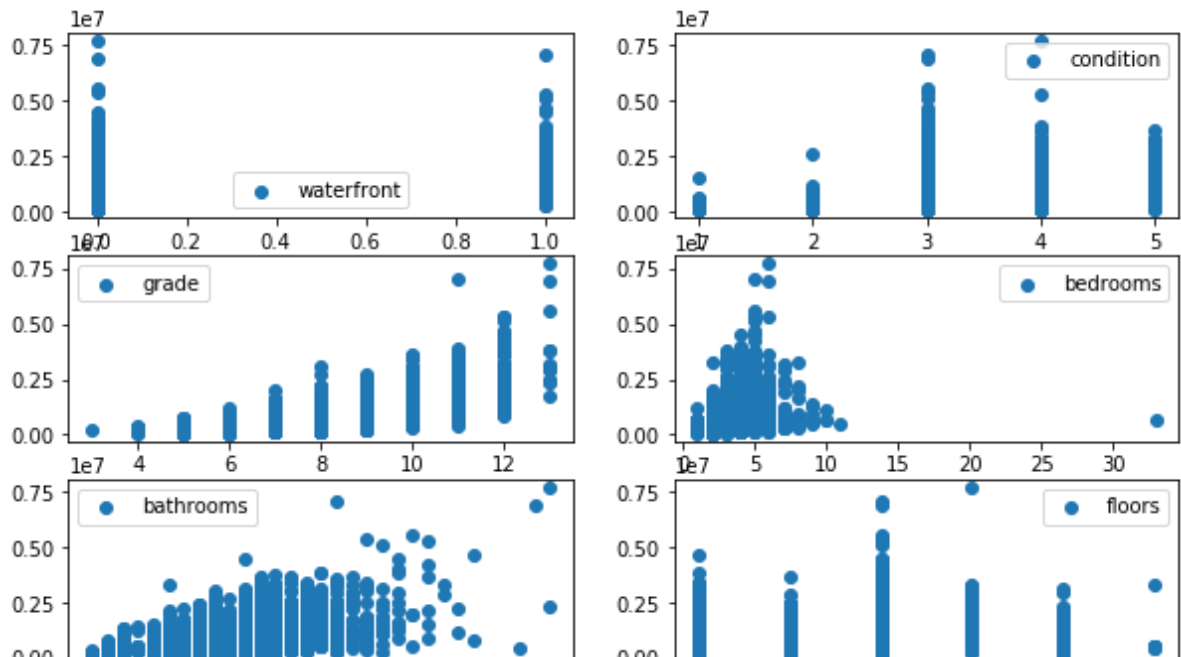
We identify that the suspected categorical variables have a limited range. So there is a good chance that they actually are.

Let's plot the price against those variables if we can see any pattern.

```
In [16]: #
import matplotlib.pyplot as plt
%matplotlib inline

num_lines = len(categorical_variables)//2 + 1
fig = plt.figure(figsize=(10, 10))

for index, xcol in enumerate(categorical_variables):
    plt.subplot(num_lines, 2, index + 1)
    plt.scatter(data=cleaned_data, x=xcol, y='price', label=xcol);
    plt.legend();
```



Conclusion : Plotting them allow us to clearly indentify categorical data and especially for the zipcode when you look at zipcode > 98150

```
In [17]: # Transform Categorical Variables using one-hot-encoding
categorical_variables = ['waterfront', 'condition', 'grade', 'bedrooms', 'bathrooms', 'floors', 'view', 'zipcode']

water_dummies = pd.get_dummies(cleaned_data['waterfront'], prefix="water")
cond_dummies = pd.get_dummies(cleaned_data['condition'], prefix="cond")
grade_dummies = pd.get_dummies(cleaned_data['grade'], prefix="grade")
bed_dummies = pd.get_dummies(cleaned_data['bedrooms'], prefix='bed')
bath_dummies = pd.get_dummies(cleaned_data['bathrooms'], prefix='bath')
floor_dummies = pd.get_dummies(cleaned_data['floors'], prefix='floor')
view_dummies = pd.get_dummies(cleaned_data['view'], prefix='view')
zipcode_dummies = pd.get_dummies(cleaned_data['zipcode'], prefix='zip')
```



```
In [18]: cleaned_data = cleaned_data.drop(categorical_variables, axis=1)

cleaned_data = pd.concat([
    cleaned_data,
    water_dummies,
    cond_dummies,
    grade_dummies,
    bed_dummies,
    bath_dummies,
    floor_dummies,
    view_dummies,
    zipcode_dummies],
    axis=1)

cleaned_data.head()
```

Out[18]:

	id	date	price	sqft_living	sqft_lot	sqft_above	sqft_basement	yr_built	yr_renovat
0	7129300520	2014-10-13	221900.0	1180	5650	1180	0.0	1955-01-01	1955-01-
1	6414100192	2014-12-09	538000.0	2570	7242	2170	400.0	1951-01-01	1991-01-
2	5631500400	2015-02-25	180000.0	770	10000	770	0.0	1933-01-01	1933-01-
3	2487200875	2014-12-09	604000.0	1960	5000	1050	910.0	1965-01-01	1965-01-
4	1954400510	2015-02-18	510000.0	1680	8080	1680	0.0	1987-01-01	1987-01-

5 rows × 153 columns

Step 4: Save our cleaned dataset for reusability

Let's save our cleaned dataset to a distinct csv file

```
In [19]: # Save clean dataset to cleaned_kc_house_data.csv file
cleaned_data.to_csv('cleaned_kc_house_data.csv', index=False)

# test loading the cleaned dataset file and see if we kept our data clean
tested_dataset = pd.read_csv('cleaned_kc_house_data.csv')
tested_dataset.head()
```

Out[19]:

	id	date	price	sqft_living	sqft_lot	sqft_above	sqft_basement	yr_built	yr_renovat
0	7129300520	2014-10-13	221900.0	1180	5650	1180	0.0	1955-01-01	1955-01-
1	6414100192	2014-12-09	538000.0	2570	7242	2170	400.0	1951-01-01	1991-01-
2	5631500400	2015-02-25	180000.0	770	10000	770	0.0	1933-01-01	1933-01-
3	2487200875	2014-12-09	604000.0	1960	5000	1050	910.0	1965-01-01	1965-01-
4	1954400510	2015-02-18	510000.0	1680	8080	1680	0.0	1987-01-01	1987-01-

5 rows × 153 columns

```
In [20]: # Check datatypes
tested_dataset.dtypes
```

```
Out[20]: id                int64
date                object
price              float64
sqft_living        int64
sqft_lot           int64
sqft_above         int64
sqft_basement      float64
yr_built           object
yr_renovated       object
lat               float64
long              float64
sqft_living15      int64
sqft_lot15         int64
water_False        int64
water_True         int64
cond_1             int64
cond_2             int64
cond_3             int64
cond_4             int64
cond_5             int64
```

```
In [21]: # We can see that the datetime format is not kept. Let's fix it
date_columns = ['date', 'yr_built', 'yr_renovated']

for date_column in date_columns:
    tested_dataset[date_column] = pd.to_datetime(tested_dataset[date_column])

tested_dataset.dtypes
```

```
Out[21]: id                int64
date                datetime64[ns]
price                float64
sqft_living          int64
sqft_lot             int64
sqft_above           int64
sqft_basement        float64
yr_built             datetime64[ns]
yr_renovated         datetime64[ns]
lat                  float64
long                 float64
sqft_living15        int64
sqft_lot15           int64
water_False          int64
water_True           int64
cond_1               int64
cond_2               int64
cond_3               int64
cond_4               int64
...                ...
```

```
In [26]: tested_dataset.columns
```

```
Out[26]: Index(['id', 'date', 'price', 'sqft_living', 'sqft_lot', 'sqft_above',
                'sqft_basement', 'yr_built', 'yr_renovated', 'lat',
                ...,
                'zip_98146', 'zip_98148', 'zip_98155', 'zip_98166', 'zip_98168',
                'zip_98177', 'zip_98178', 'zip_98188', 'zip_98198', 'zip_98199'],
                dtype='object', length=153)
```

Conclusion

All looks good. We managed to clean our data and implement a way to reuse it once we will work on regression.